

Smart Financial Coach – Design Documentation

Smart Financial Coach – Design Documentation

1. App Design Overview

The **Smart Financial Coach** is a local-first financial assistant that combines **machine learning regression** with a **generative AI orchestrator**.

- **User Interaction:** Command-line interface (CLI).
- **Data Input:** User provides a `.csv` file of bank transactions (`Date`, `Description`, `Amount`, `Running Bal.`).
- **Orchestrator (LLM):** A local LLM (via Ollama) interprets user queries, classifies them into one of three categories:
 - **Regression task:** Forecast future balance using ML.
 - **General task:** Provide advice or explanations using the LLM.
 - **Clarification task:** Ask the user for missing parameters (e.g., time horizon).
- **ML Forecasting Engine:**
 - Uses **Polynomial Regression** (degrees 1–3) with **cross-validation** to forecast account balances.
 - Handles both relative queries (“in 30 days”) and absolute queries (“on August 10th”).
- **Session Memory:** Keeps past conversation context to support follow-up questions.
- **Responsible AI:**
 - Data stays local; no external services.
 - Deterministic ML for numeric forecasts, AI only for interpretation and explanations.

2. Tech Stack

- **Programming Language:** Python 3.9+
- **Libraries:**
 - `requests` → Communicate with Ollama API
 - `numpy` → Numeric operations
 - `scikit-learn` → Polynomial regression (`LinearRegression` + `PolynomialFeatures`)
 - `matplotlib` (optional) → Visualizations (used in presentation/demo)
- **LLM Runtime:** [Ollama](https://ollama.ai/) for running local models (e.g., `llama3.1`)
- **Interface:** Command-line interface (CLI)
- **Data Format:** CSV (comma-separated values, with balances and transactions)

3. Design Components

a. CLI Layer

- Handles user input/output
- Manages session loop (exit/quit commands supported)

b. Orchestrator (LLM via Ollama)

- Classifies queries into regression/general/clarify
- Parses absolute dates (` August 10th`) and converts them into ISO format
- Extracts numeric parameters (e.g., horizon_days, target thresholds)

c. Regression Forecasting Engine

- Builds a **time series** of balances from CSV
- Fits **Polynomial Regression** (degree auto-selected via CV)
- Outputs forecast date, balance, model degree, and cross-validation RMSE

d. Session Memory

- Stores prior user and assistant messages
- Provides history to the orchestrator for contextual understanding

4. Example User Flows

1. **Forecasting Query (Regression)**

- User: **"How much money will I have in my account on August 10th?"**
- Orchestrator → classify as regression, extract target date.
- Forecasting engine → predict balance on 2025-08-10.
- Assistant: **"My ML regression forecast for 2025-08-10 is \$X (degree=2, CV_RMSE=Y)."**

2. **Advice Query (General)**

- User: **"How can I cut down on my spending?"**
- Orchestrator → classify as general.
- Assistant: **"You could reduce dining out expenses and set a monthly budget limit."**

3. **Clarification Query**

- User: **"Will I have enough money soon?"**
- Orchestrator → classify as clarify.
- Assistant: **"Could you specify how many days or weeks ahead you want me to forecast?"**

5. Future Enhancements

1. **UI Layer:** Add a simple web or mobile interface.
2. **Visual Forecasts:** Display balance projections as interactive charts.
3. **Advanced ML Models:** Incorporate ARIMA or LSTM for time-series forecasting.
4. **Subscription Detection:** Implement pattern recognition for recurring charges.
5. **Budgeting Features:** Allow goal-setting and alerts for overspending.
6. **External API Integration:** Connect to real-time bank feeds (with secure tokens).
7. **Personalization:** Provide custom saving strategies based on spending habits.

8. **Privacy Enhancements:** Explore federated learning to improve models across users without sharing raw data.